

```
In [ ]: # This Python 3 environment comes with many helpful analytics libraries i
# It is defined by the kaggle/python Docker image: https://github.com/kag
# For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) wil

import os
# for dirname, _, filenames in os.walk('/kaggle/input'):
#     for filename in filenames:
#         print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/) th
# You can also write temporary files to /kaggle/temp/, but they won't be
```

```
In [ ]: import numpy as np
import matplotlib.pyplot as plt
import torch
import torch as th
from torch import nn
import torchvision
import torchvision.transforms as transforms
from torchvision.datasets import ImageFolder
import torchvision.models as models
from torch.utils.data import TensorDataset, ConcatDataset, DataLoader, Sub
from sklearn.model_selection import KFold
import os
from torchinfo import summary
```

```
In [ ]: device = "cuda" if torch.cuda.is_available() else "cpu"
```

```
In [ ]: # ResNet 10 Approach
```

```
In [6]: import timm
class Resnet10(nn.Module):
    def __init__(self, num_classes=3):
        super(Resnet10, self).__init__()
        self.resnet10 = timm.create_model("resnet10t", pretrained=True)

        # Replace the fully connected layer for classification
        in_features = self.resnet10.get_classifier().in_features
        self.resnet10.fc = nn.Linear(in_features, num_classes)

    def forward(self, x):
        return self.resnet10(x)
```

```
In [8]: torch.cuda.empty_cache()
```

```
In [9]: encoder = Resnet10(num_classes=3).to(device)
```

```
In [13]: transformer=transforms.Compose([
    transforms.Resize((150,150)),
    transforms.RandomHorizontalFlip(),
```

```

        transforms.RandomRotation(degrees=10),
        transforms.ToTensor(), #0-255 to 0-1, numpy to tensors
        transforms.Normalize([0.5,0.5,0.5], # 0-1 to [-1,1] , formula (x-mean
                                [0.5,0.5,0.5])
    ])

```

```

In [14]: # Mentioning training path
train_path = '/kaggle/input/gsoc123/new_dataset/train'

```

```

In [15]: class_folders = torchvision.datasets.ImageFolder(train_path,transform=tra

```

```

In [16]: from torch.optim.lr_scheduler import Steplr
# Define the scheduler

```

```

In [17]: def CNN_train(epochs,class_folders, num_folds, lr):
    loss_fn = nn.CrossEntropyLoss()
    train_acc_values = []
    best_accuracy = 0.0
    test_acc_values = []
    num_epochs = epochs
    epoch_count = []
    iteration_details = []
    optimizer = torch.optim.Adam(encoder.parameters(), lr=lr, weight_decay=1e-4)
    scheduler = Steplr(optimizer, step_size=10, gamma=0.1, verbose=True)

    # Initialize KFold object
    kf = KFold(n_splits=num_folds, shuffle=True)

    # Loop through each fold
    for fold, (train_index, val_index) in enumerate(kf.split(class_folders)):
        print(f"Fold {fold + 1}/{num_folds}")

        dataset_train = Subset(class_folders, train_index)
        dataset_valid = Subset(class_folders, val_index)

        train_loader = torch.utils.data.DataLoader(
            dataset_train, batch_size=64, shuffle=True
        )
        val_loader = torch.utils.data.DataLoader(
            dataset_valid, batch_size=32, shuffle=True
        )

        # Training loop
        for epoch in range(num_epochs):
            # Set model to training mode
            encoder.train()

            train_accuracy = 0.0
            train_loss = 0.0

            for i, (images, labels) in enumerate(train_loader):
                if torch.cuda.is_available():
                    images = images.cuda()
                    labels = labels.cuda()

                optimizer.zero_grad()
                outputs = encoder(images)
                loss = loss_fn(outputs, labels)
                loss.backward()

```

```

optimizer.step()

train_loss += loss.item() * images.size(0)
_, prediction = torch.max(outputs.data, 1)
train_accuracy += int(torch.sum(prediction == labels.data

train_accuracy = train_accuracy / len(train_index)
train_loss = train_loss / len(train_index)

# Validation loop
encoder.eval()

val_accuracy = 0.0
val_loss = 0.0

for i, (images, labels) in enumerate(val_loader):
    if torch.cuda.is_available():
        images = images.cuda()
        labels = labels.cuda()

    outputs = encoder(images)
    loss = loss_fn(outputs, labels)
    val_loss += loss.item()* images.size(0) # Accumulate the
    _, prediction = torch.max(outputs.data, 1)
    val_accuracy += int(torch.sum(prediction == labels.data))

# Compute average loss and accuracy
val_loss /= len(val_index)
val_accuracy = val_accuracy / len(val_index)

# Step the scheduler
scheduler.step()

print(f"Epoch [{epoch + 1}/{num_epochs}], Train Loss: {train_

# Save the best model
if val_accuracy > best_accuracy:
    torch.save(encoder.state_dict(), 'best_model.pth')
    best_accuracy = val_accuracy

print(f"Best Validation Accuracy: {best_accuracy}")

```

In [18]: results =CNN_train(3,class_folders,10,0.001)

```

/usr/local/lib/python3.10/dist-packages/torch/optim/lr_scheduler.py:62: Us
erWarning: The verbose parameter is deprecated. Please use get_last_lr() t
o access the learning rate.
warnings.warn(

```

Fold 1/10

Epoch [1/3], Train Loss: 1.094, Train Accuracy: 0.361, Val Loss: 1.136, Validation Accuracy: 0.355

Epoch [2/3], Train Loss: 0.966, Train Accuracy: 0.496, Val Loss: 1.014, Validation Accuracy: 0.507

Epoch [3/3], Train Loss: 0.721, Train Accuracy: 0.674, Val Loss: 0.586, Validation Accuracy: 0.758

Fold 2/10

Epoch [1/3], Train Loss: 0.526, Train Accuracy: 0.783, Val Loss: 0.492, Validation Accuracy: 0.802

Epoch [2/3], Train Loss: 0.445, Train Accuracy: 0.823, Val Loss: 0.479, Validation Accuracy: 0.829

Epoch [3/3], Train Loss: 0.403, Train Accuracy: 0.843, Val Loss: 0.521, Validation Accuracy: 0.782

Fold 3/10

Epoch [1/3], Train Loss: 0.366, Train Accuracy: 0.857, Val Loss: 0.473, Validation Accuracy: 0.837

Epoch [2/3], Train Loss: 0.350, Train Accuracy: 0.866, Val Loss: 0.365, Validation Accuracy: 0.862

Epoch [3/3], Train Loss: 0.327, Train Accuracy: 0.875, Val Loss: 0.454, Validation Accuracy: 0.814

Fold 4/10

Epoch [1/3], Train Loss: 0.319, Train Accuracy: 0.879, Val Loss: 0.391, Validation Accuracy: 0.867

Epoch [2/3], Train Loss: 0.244, Train Accuracy: 0.909, Val Loss: 0.217, Validation Accuracy: 0.921

Epoch [3/3], Train Loss: 0.228, Train Accuracy: 0.917, Val Loss: 0.203, Validation Accuracy: 0.926

Fold 5/10

Epoch [1/3], Train Loss: 0.221, Train Accuracy: 0.918, Val Loss: 0.213, Validation Accuracy: 0.922

Epoch [2/3], Train Loss: 0.212, Train Accuracy: 0.922, Val Loss: 0.198, Validation Accuracy: 0.923

Epoch [3/3], Train Loss: 0.209, Train Accuracy: 0.924, Val Loss: 0.209, Validation Accuracy: 0.921

Fold 6/10

Epoch [1/3], Train Loss: 0.206, Train Accuracy: 0.923, Val Loss: 0.176, Validation Accuracy: 0.933

Epoch [2/3], Train Loss: 0.196, Train Accuracy: 0.927, Val Loss: 0.180, Validation Accuracy: 0.935

Epoch [3/3], Train Loss: 0.195, Train Accuracy: 0.928, Val Loss: 0.177, Validation Accuracy: 0.939

Fold 7/10

Epoch [1/3], Train Loss: 0.188, Train Accuracy: 0.931, Val Loss: 0.177, Validation Accuracy: 0.935

Epoch [2/3], Train Loss: 0.183, Train Accuracy: 0.934, Val Loss: 0.170, Validation Accuracy: 0.934

Epoch [3/3], Train Loss: 0.175, Train Accuracy: 0.935, Val Loss: 0.167, Validation Accuracy: 0.938

Fold 8/10

Epoch [1/3], Train Loss: 0.170, Train Accuracy: 0.938, Val Loss: 0.160, Validation Accuracy: 0.944

Epoch [2/3], Train Loss: 0.169, Train Accuracy: 0.938, Val Loss: 0.166, Validation Accuracy: 0.938

Epoch [3/3], Train Loss: 0.172, Train Accuracy: 0.937, Val Loss: 0.162, Validation Accuracy: 0.939

Fold 9/10

Epoch [1/3], Train Loss: 0.169, Train Accuracy: 0.939, Val Loss: 0.161, Validation Accuracy: 0.945

Epoch [2/3], Train Loss: 0.165, Train Accuracy: 0.939, Val Loss: 0.164, Va

Validation Accuracy: 0.939
 Epoch [3/3], Train Loss: 0.168, Train Accuracy: 0.938, Val Loss: 0.163, Validation Accuracy: 0.936
 Fold 10/10
 Epoch [1/3], Train Loss: 0.163, Train Accuracy: 0.941, Val Loss: 0.160, Validation Accuracy: 0.939
 Epoch [2/3], Train Loss: 0.163, Train Accuracy: 0.941, Val Loss: 0.154, Validation Accuracy: 0.940
 Epoch [3/3], Train Loss: 0.161, Train Accuracy: 0.941, Val Loss: 0.167, Validation Accuracy: 0.936
 Best Validation Accuracy: 0.9446666666666667

```

In [20]: test_path='/kaggle/input/gsoc123/new_dataset/val'
test_loader=DataLoader(
    torchvision.datasets.ImageFolder(test_path,transform=transformer),
    batch_size=32, shuffle=True
)
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import roc_curve, auc
from sklearn.preprocessing import label_binarize

# Assuming model is already defined and moved to GPU if available

# Assuming transformer is defined

# Assuming test_loader is defined

y_score_list = []
y_true_list = []

# Evaluate model
encoder.eval()
for images, labels in test_loader:
    if torch.cuda.is_available():
        images = images.cuda()
        labels = labels.cuda()

    with torch.no_grad():
        y_score_batch = encoder(images)

    y_score_list.append(y_score_batch.cpu().numpy())
    y_true_list.append(labels.cpu().numpy())

y_score = np.vstack(y_score_list)
y_true = np.hstack(y_true_list)

# Binarize the ground truth labels
y_true_bin = label_binarize(y_true, classes=np.unique(y_true))

# Compute ROC curve and ROC area for each class
n_classes = y_score.shape[1]
fpr = dict()
tpr = dict()
roc_auc = dict()

for i in range(n_classes):
    fpr[i], tpr[i], _ = roc_curve(y_true_bin[:, i], y_score[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])
  
```

```

# Compute micro-average ROC curve and ROC area
fpr["micro"], tpr["micro"], _ = roc_curve(y_true_bin.ravel(), y_score.ravel())
roc_auc["micro"] = auc(fpr["micro"], tpr["micro"])

# Plot ROC curve
plt.figure()
plt.plot(fpr["micro"], tpr["micro"], color='deeppink', lw=2, label=f'ROC curve (micro-AUC = {roc_auc["micro"]:.2f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve (Multiclass)')
plt.legend(loc='lower right')
plt.show()

```

