

```
In [ ]: import numpy as np
import matplotlib.pyplot as plt
import torch
from torch import nn
import torchvision
import torchvision.transforms as transforms
from torchvision.datasets import ImageFolder
import torchvision.models as models
from torch.utils.data import TensorDataset, ConcatDataset, DataLoader, Sub
from sklearn.model_selection import KFold
import os
from torchinfo import summary
```

```
In [4]: device = "cuda" if torch.cuda.is_available() else "cpu"
```

```
In [ ]: class DeepLens(nn.Module):
    def __init__(self, normalized_r2):
        super(DeepLens, self).__init__()
        self.normalized_r2 = normalized_r2

        # Step 2: Use another ResNet-18 for vector A
        self.resnet_A = models.resnet18(pretrained=True)

        #         for param in self.resnet_B.parameters(): # Freeze ResNet-18 pa
        #             param.requires_grad = False

        self.features_A = nn.Sequential(*list(self.resnet_A.children())[:
self.features_A.add_module('flatten', nn.Flatten())
self.features_A.add_module('dropout', nn.Dropout(p=0.4)) # Addin
self.relu = nn.ReLU(inplace=True)
self.features_A.add_module('linear', nn.Linear(12800, 22500))

        # Step 3: Use another ResNet-18 for vector B
        self.resnet_B = models.resnet18(pretrained=True)

        #         for param in self.resnet_B.parameters(): # Freeze ResNet-18 pa
        #             param.requires_grad = False

        self.features_B = nn.Sequential(*list(self.resnet_B.children())[:
self.features_B.add_module('flatten', nn.Flatten())
self.features_B.add_module('dropout', nn.Dropout(p=0.4)) # Addin
#         self.relu = nn.ReLU(inplace=True)
self.features_B.add_module('linear', nn.Linear(12800, 22500))

        # Step 5: Neural layer and softmax

        self.neural_layer1 = nn.Linear(22500, 5625 )
        self.neural_layer2 = nn.Linear(5625, 128)
        self.neural_layer3 = nn.Linear(128, 3)
        self.dropout = nn.Dropout(p=0.4) # Adding dropout layer

    def forward(self, x):

        # Step 2: Extract features A

        features_A = self.features_A(x)
        features_A = self.relu(features_A)
        features_A = features_A.view(features_A.size(0), -1)
```

```

# Step 3: Extract features B

features_B = self.features_B(x)
features_B = self.relu(features_B)
features_B = features_B.view(features_B.size(0), -1)

# Step 4: Physics equation  $C = I - B * (\text{distance from center of } i)$ 
C = features_A - features_B * (self.normalized_r2)

# Step 5: Neural layer and softmax
output = self.neural_layer1(C)
output = self.relu(output)
output = self.dropout(output)

output = self.neural_layer2(output)
output = self.relu(output)
output = self.dropout(output) # Applying dropout

output = self.neural_layer3(output)

return output

```

```

In [7]: transformer=transforms.Compose([
    transforms.Resize((150,150)),
    transforms.RandomHorizontalFlip(),
    transforms.RandomRotation(degrees=10),
    transforms.ToTensor(), #0-255 to 0-1, numpy to tensors
    transforms.Normalize([0.5,0.5,0.5], # 0-1 to [-1,1] , formula (x-mean
        [0.5,0.5,0.5])
])

```

```

In [8]: grid_size = 150

# Initialize an empty tensor to store distances
r2 = torch.zeros(grid_size * grid_size)

# Compute distances for each pixel
for i in range(grid_size):
    for j in range(grid_size):
        r2[grid_size * i + j] = ((74.5 - i) ** 2 + (74.5 - j) ** 2)

# Verify the size of dis
print("Size of dis tensor:", r2.size())

# Find the maximum value in the r2 tensor
max_r2 = r2.max()
normalized_r2 = r2/max_r2

# Move r2 tensor to CUDA if available
normalized_r2 = normalized_r2.cuda()

```

Size of dis tensor: torch.Size([22500])

```

In [9]: encoder = DeepLens(normalized_r2).to(device)

```

```

/usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:208:
UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be
removed in the future, please use 'weights' instead.
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:223:
UserWarning: Arguments other than a weight enum or `None` for 'weights' ar
e deprecated since 0.13 and may be removed in the future. The current beha
vior is equivalent to passing `weights=ResNet18_Weights.IMAGENET1K_V1`. Yo
u can also use `weights=ResNet18_Weights.DEFAULT` to get the most up-to-da
te weights.
  warnings.warn(msg)
Downloading: "https://download.pytorch.org/models/resnet18-f37072fd.pth" to
/root/.cache/torch/hub/checkpoints/resnet18-f37072fd.pth
100%|██████████| 44.7M/44.7M [00:00<00:00, 148MB/s]

```

```
In [10]: train_path = '/kaggle/input/gsoc-pinn1/new_dataset_PINN/train'
```

```
In [11]: class_folders = torchvision.datasets.ImageFolder(train_path, transform=tra
```

```
In [12]: torch.cuda.empty_cache()
```

```
In [13]: from torch.optim.lr_scheduler import StepLR
optimizer = torch.optim.RMSprop(encoder.parameters(), lr=0.01, weight_dec
# Define the scheduler
scheduler = StepLR(optimizer, step_size=10, gamma=0.1, verbose=True)
```

```

/usr/local/lib/python3.10/dist-packages/torch/optim/lr_scheduler.py:62: Us
erWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
  warnings.warn(

```

```
In [16]: num_epochs = 10
num_folds = 3
best_accuracy = 0.0
loss_fn = nn.CrossEntropyLoss()
# Initialize KFold object
kf = KFold(n_splits=num_folds, shuffle=True)

# Loop through each fold
for fold, (train_index, val_index) in enumerate(kf.split(class_folders)):
    print(f"Fold {fold + 1}/{num_folds}")

    dataset_train = Subset(class_folders, train_index)
    dataset_valid = Subset(class_folders, val_index)

    train_loader = torch.utils.data.DataLoader(
        dataset_train, batch_size=64, shuffle=True
    )
    val_loader = torch.utils.data.DataLoader(
        dataset_valid, batch_size=32, shuffle=True
    )

    # Training loop
    for epoch in range(num_epochs):
        # Set model to training mode
        encoder.train()

        train_accuracy = 0.0
        train_loss = 0.0

```

```

for i, (images, labels) in enumerate(train_loader):
    if torch.cuda.is_available():
        images = images.cuda()
        labels = labels.cuda()

    optimizer.zero_grad()
    outputs = encoder(images)
    loss = loss_fn(outputs, labels)
    loss.backward()
    optimizer.step()

    train_loss += loss.item() * images.size(0)
    _, prediction = torch.max(outputs.data, 1)
    train_accuracy += int(torch.sum(prediction == labels.data))

train_accuracy = train_accuracy / len(train_index)
train_loss = train_loss / len(train_index)

# Validation loop
encoder.eval()

val_accuracy = 0.0
val_loss = 0.0

for i, (images, labels) in enumerate(val_loader):
    if torch.cuda.is_available():
        images = images.cuda()
        labels = labels.cuda()

    outputs = encoder(images)
    loss = loss_fn(outputs, labels)
    val_loss += loss.item() * images.size(0) # Accumulate the loss
    _, prediction = torch.max(outputs.data, 1)
    val_accuracy += int(torch.sum(prediction == labels.data))

# Compute average loss and accuracy
val_loss /= len(val_index)
val_accuracy = val_accuracy / len(val_index)

print(f"Epoch [{epoch + 1}/{num_epochs}], Train Loss: {train_loss}")

# Save the best model
if val_accuracy > best_accuracy:
    torch.save(encoder.state_dict(), 'best_model.pth')
    best_accuracy = val_accuracy

print(f"Best Validation Accuracy: {best_accuracy}")

```

Fold 1/5
Epoch [1/5], Train Loss: 1.101, Train Accuracy: 0.340, Val Loss: 1.096, Validation Accuracy: 0.368
Epoch [2/5], Train Loss: 1.096, Train Accuracy: 0.351, Val Loss: 1.086, Validation Accuracy: 0.387
Epoch [3/5], Train Loss: 1.053, Train Accuracy: 0.419, Val Loss: 1.004, Validation Accuracy: 0.459
Epoch [4/5], Train Loss: 0.902, Train Accuracy: 0.547, Val Loss: 0.821, Validation Accuracy: 0.614
Epoch [5/5], Train Loss: 0.734, Train Accuracy: 0.669, Val Loss: 0.759, Validation Accuracy: 0.648
Fold 2/5
Epoch [1/5], Train Loss: 0.614, Train Accuracy: 0.740, Val Loss: 0.893, Validation Accuracy: 0.647
Epoch [2/5], Train Loss: 0.538, Train Accuracy: 0.779, Val Loss: 0.933, Validation Accuracy: 0.626
Epoch [3/5], Train Loss: 0.493, Train Accuracy: 0.800, Val Loss: 0.484, Validation Accuracy: 0.804
Epoch [4/5], Train Loss: 0.457, Train Accuracy: 0.816, Val Loss: 0.515, Validation Accuracy: 0.799
Epoch [5/5], Train Loss: 0.421, Train Accuracy: 0.833, Val Loss: 0.471, Validation Accuracy: 0.813
Fold 3/5
Epoch [1/5], Train Loss: 0.410, Train Accuracy: 0.839, Val Loss: 0.924, Validation Accuracy: 0.737
Epoch [2/5], Train Loss: 0.382, Train Accuracy: 0.852, Val Loss: 0.548, Validation Accuracy: 0.802
Epoch [3/5], Train Loss: 0.367, Train Accuracy: 0.858, Val Loss: 0.504, Validation Accuracy: 0.826
Epoch [4/5], Train Loss: 0.345, Train Accuracy: 0.866, Val Loss: 0.406, Validation Accuracy: 0.842
Epoch [5/5], Train Loss: 0.336, Train Accuracy: 0.871, Val Loss: 0.557, Validation Accuracy: 0.775
Fold 4/5
Epoch [1/5], Train Loss: 0.335, Train Accuracy: 0.872, Val Loss: 0.338, Validation Accuracy: 0.866
Epoch [2/5], Train Loss: 0.317, Train Accuracy: 0.879, Val Loss: 0.378, Validation Accuracy: 0.859
Epoch [3/5], Train Loss: 0.310, Train Accuracy: 0.883, Val Loss: 0.353, Validation Accuracy: 0.858
Epoch [4/5], Train Loss: 0.297, Train Accuracy: 0.887, Val Loss: 0.299, Validation Accuracy: 0.882
Epoch [5/5], Train Loss: 0.291, Train Accuracy: 0.889, Val Loss: 0.298, Validation Accuracy: 0.885
Fold 5/5
Epoch [1/5], Train Loss: 0.285, Train Accuracy: 0.892, Val Loss: 0.299, Validation Accuracy: 0.888
Epoch [2/5], Train Loss: 0.278, Train Accuracy: 0.896, Val Loss: 0.334, Validation Accuracy: 0.877
Epoch [3/5], Train Loss: 0.270, Train Accuracy: 0.899, Val Loss: 0.387, Validation Accuracy: 0.869
Epoch [4/5], Train Loss: 0.258, Train Accuracy: 0.904, Val Loss: 0.301, Validation Accuracy: 0.891
Epoch [5/5], Train Loss: 0.258, Train Accuracy: 0.904, Val Loss: 0.266, Validation Accuracy: 0.900
Best Validation Accuracy: 0.9003333333333333

```
In [19]: test_path='/kaggle/input/gsoc-pinn1/new_dataset_PINN/val'
test_loader=DataLoader(
    torchvision.datasets.ImageFolder(test_path,transform=transformer),
```

```

    batch_size=32, shuffle=True
)
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import roc_curve, auc
from sklearn.preprocessing import label_binarize

y_score_list = []
y_true_list = []

# Evaluate model
encoder.eval()
for images, labels in test_loader:
    if torch.cuda.is_available():
        images = images.cuda()
        labels = labels.cuda()

    with torch.no_grad():
        y_score_batch = encoder(images)

    y_score_list.append(y_score_batch.cpu().numpy())
    y_true_list.append(labels.cpu().numpy())

y_score = np.vstack(y_score_list)
y_true = np.hstack(y_true_list)

# Binarize the ground truth labels
y_true_bin = label_binarize(y_true, classes=np.unique(y_true))

# Compute ROC curve and ROC area for each class
n_classes = y_score.shape[1]
fpr = dict()
tpr = dict()
roc_auc = dict()

for i in range(n_classes):
    fpr[i], tpr[i], _ = roc_curve(y_true_bin[:, i], y_score[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

# Compute micro-average ROC curve and ROC area
fpr["micro"], tpr["micro"], _ = roc_curve(y_true_bin.ravel(), y_score.ravel())
roc_auc["micro"] = auc(fpr["micro"], tpr["micro"])

# Plot ROC curve
plt.figure()
plt.plot(fpr["micro"], tpr["micro"], color='deeppink', lw=2, label=f'ROC')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve (Multiclass)')
plt.legend(loc='lower right')
plt.show()

```

Receiver Operating Characteristic (ROC) Curve (Multiclass)

